

---

# Cascaded Transfer: Learning Many Tasks under Budget Constraints

---

Anonymous Authors<sup>1</sup>

## Abstract

*Many-Task Learning* refers to the setting where a large number of related tasks need to be learned, the exact relationships between tasks are not known, and budget constraints are in place. We introduce the *Cascaded Transfer Learning*, a novel many-task transfer learning paradigm where information (e.g. model parameters) cascades hierarchically through tasks that are learned by individual models of the same class, while respecting given budget constraints. The cascade is organized as a rooted tree that specifies the order in which tasks are learned and refined. We design a cascaded transfer mechanism deployed over a minimum spanning tree structure that connects the tasks according to a suitable distance measure, and allocates the available training budget along its branches. Experiments on synthetic and real many-task settings show that the resulting method enables more accurate and cost-effective adaptation across large task collections compared to alternative approaches.

## 1. Introduction

Modern learning systems increasingly operate in settings where a large number of related tasks must be handled under computational or training-budget constraints. This arises in domains involving personalized, localized, or fine-grained problems, where each task comes with limited data and solving all tasks independently becomes inefficient. Exploiting relationships across tasks is essential for achieving accurate models while maintaining feasible computational cost in such a context.

*Multi-Task Learning* (MTL) has been specifically developed for leveraging shared information across tasks, seeking to improve performance by joint training of multiple related tasks. The main MTL method families include feature-

sharing architectures, low-rank parameterizations, task clustering and other methods that learn task relationships (Zhang & Yang, 2018; 2021; Ruiz et al., 2024). Although effective, they typically require synchronized and globally coordinated optimization for training across all tasks, which lead to substantial memory or communication overhead. Moreover, relying on accurate task relations is a sensitive point for MTL, hence there are methods that try to infer those relations, but in the literature this is most usually required to be apriori known.

As the number of tasks grows, these limitations become more pronounced, giving rise to the *Many-Task Learning* (MaTL) setting. Here, the focus shifts to handling hundreds or thousands of related tasks, whose relatedness is typically unknown or only partially observable, and often under resource constraints that make joint optimization impractical. The MaTL regime appears naturally in personalized modeling, distributed sensing, or large-scale prediction environments, where one seeks scalable mechanisms for exploiting task relatedness. Energy networks, transportation, retail, where thousands of related tasks arise in parallel (He et al., 2019), are among the typical application sectors that fall within the MaTL regime. Prior work has shown that scaling to many tasks requires explicit mechanisms to control parameter sharing and interference within a *single* model, such as task-specific routing (Strezoski et al., 2019), hierarchical many-task architectures (Hashimoto et al., 2017; Liu et al., 2021), or transferability-based task selection that guides efficient knowledge reuse across large task collections (Tan et al., 2024). Studies of extreme multi-task pre-training with over one hundred tasks further highlight the benefits and the complexity of this regime (Aribandi et al., 2022). In parallel, foundation models have recently been explored for graphs, with the goal of capturing generalities across heterogeneous graph tasks. A particularly illustrative instance of these pressures is found in modern time-series forecasting. Transformer-based architectures such as PatchTST (Nie et al., 2022) and universal forecasters like Moirai (Woo et al., 2024) demonstrate that shared representations can be remarkably effective across a wide variety of series and domains. Yet these approaches typically rely on large-scale centralized pretraining and assume that inference or finetuning can be performed without stringent per-task budget constraints. In applications involving thousands of forecast-

<sup>1</sup>Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

055  
056  
057  
058  
059  
060  
061  
062  
063  
064  
065  
066  
067  
068  
069  
070  
071  
072  
073  
074  
075  
076  
077  
078  
079  
080  
081  
082  
083  
084  
085  
086  
087  
088  
089  
090  
091  
092  
093  
094  
095  
096  
097  
098  
099  
100  
101  
102  
103  
104  
105  
106  
107  
108  
109

ing tasks (each requiring lightweight, task-specific adaptation) global training becomes costly or infeasible, and the problem shifts toward designing scalable mechanisms for transferring information across many related tasks.

Thinking about sharing information from one task to assist another without requiring global joint training, *Transfer Learning* (TL) comes naturally into the discussion. TL methods are categorized to instance-based, feature-based, parameter-based and relational transfer (Pan & Yang, 2009; Weiss et al., 2016; Zhuang et al., 2020). In practice, TL can be straightforward to deploy in modern neural settings and scales well to new tasks, but it treats each source–target pair independently. This pairwise structure limits its applicability into the MaTL regime, as it provides only simple deployment schemes such as star-shaped transfers from a single source model. Beyond such configurations, classical TL offers little guidance on how to organize or coordinate transfers across a large collection of tasks, which makes it inadequate as the primary mechanism for structuring information flow in MTL, not to mention the MaTL regime. We focus on a frugal many-task regime in which each task is refined only once under a strict global budget, and propose *Cascaded Transfer Learning* (CTL) as a single-pass alternative to joint multi-task optimization when repeated task revisiting is infeasible.

**Contributions.** In this work, we propose a new approach that occupies the conceptual space between these two paradigms, MTL and TL. Rather than adapting a source model directly to each target task, we introduce *Cascaded Transfer Learning*, a mechanism that propagates knowledge through a sequence of tasks (see Figure 1). The key idea is that coordinated transfer across a graph structure connecting the tasks can realize gains that are not accessible through isolated, direct adaptations. To the best of our knowledge, this is the first work to provide a theoretical and algorithmic framework for budget-constrained cascaded transfer over large task collections. Our contributions are the following:

- (i) We provide a theoretical analysis of cascaded transfer over trees, establishing sufficient conditions under which CTL provably improves over direct transfer. Our results characterize how the stability of local optimization, task similarity, and cascade depth jointly control error propagation, and when tree-structured transfer is beneficial.
- (ii) Building on this analysis, we propose a scalable CTL algorithm that constructs a transfer tree and allocates a limited training budget across tasks. We empirically validate the approach on synthetic benchmarks, multisite forecasting, and image classification, showing consistent improvements over independent training and standard transfer baselines under equal or lower computational budgets.

**Notations.** Let  $\mathcal{V}$  denote the set of tasks, with  $|\mathcal{V}|$  its cardinality. For two tasks  $u, v \in \mathcal{V}$ , we write  $u \rightarrow v$  to denote a transfer from task  $u$  to task  $v$ . The norms  $\|\cdot\|$  and  $\|\cdot\|_F$  respectively denote the Euclidean norm and the Frobenius norm. For a symmetric matrix  $\mathbf{A}$ , the notation  $\mathbf{A} \succ 0$  indicates that  $\mathbf{A}$  is positive definite. Finally,  $\lceil \cdot \rceil$  denotes the ceiling function.

## 2. Preliminaries

For each task  $v \in \mathcal{V}$ , let  $\mathbf{X}_v \in \mathbb{R}^{n \times d}$  denote the task-specific data matrix, where  $n \in \mathbb{N}^*$  is the number of samples available for task  $v$  and  $d \in \mathbb{N}^*$  is the feature dimension. Each task is associated with an unknown parameter vector  $\theta_v \in \mathbb{R}^d$  to be learned from data. For each task  $v$ , let  $b_v \in \mathbb{N}^*$  denote the computational budget allocated to that task. Throughout this work, a *budget* denotes a fixed amount of local optimization effort, measured in refinement steps. We denote by  $G_v^{b_v} : \mathbb{R}^d \rightarrow \mathbb{R}^d$  a refinement operator corresponding to  $b_v$  iterations of a gradient-based optimization algorithm applied to the task-specific parameter. Given a source task  $u$  and a target task  $v$ , a transfer  $u \rightarrow v$  consists of initializing the parameters of task  $v$  using information from task  $u$ , followed by local refinement via  $G_v^{b_v}$ . This source–target transfer constitutes the elementary building block of cascaded transfer learning. To coordinate multiple such transfers across tasks, we introduce graph-structured dependency relations over  $\mathcal{V}$ , which specify how knowledge is propagated and refined throughout the task set.

In this work, we focus on *rooted trees* as the primary structure for organizing transfer. A rooted tree is a *directed acyclic graph* (DAG) in which each node has at most one parent, denoted  $\text{pa}(v)$ . The unique node with zero in-degree is called the *root*, while nodes with zero out-degree are referred to as *leaves*. This structure induces a natural flow of information from the root toward the leaves, defining an unambiguous order in which tasks are processed. More generally, a DAG is any directed graph without cycles and admits a topological ordering such that, for each edge  $(u, v)$ , node  $u$  precedes node  $v$  in the order. In this broader setting, a task may receive information from multiple parent tasks, which can be combined through an aggregation operator. While this abstraction is useful for unifying different transfer mechanisms, both our theoretical analysis and algorithmic instantiations are restricted to rooted trees, which capture the essential behavior of cascaded transfer while remaining analytically and computationally tractable.

### 2.1. The Cascaded Transfer Learning Paradigm

We now formalize CTL, starting from the rooted tree setting.

**Definition 2.1** (Cascaded Transfer Learning). Let  $\mathcal{T} = (\mathcal{V}, \mathcal{E})$  be a rooted tree whose nodes correspond to learning tasks. *Cascaded Transfer Learning* is a learning process

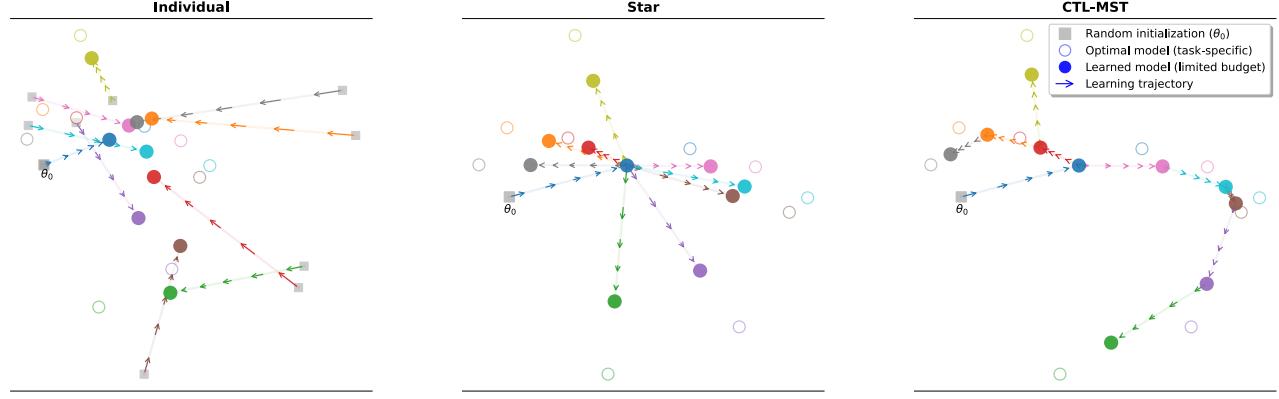


Figure 1. Parameter-space intuition for CTL. Left: Independent training. Each task is optimized from its own initialization and stops short of its optimum due to limited budget. Middle: Star transfer. One source task is learned first and directly transferred to all other tasks. Right: CTL with an MST. Tasks are learned sequentially along a minimum spanning tree, decomposing long transfers into local steps. Solid nodes denote task optima, white nodes learned parameters, and arrows optimization trajectories.

### Algorithm 1 Cascaded Transfer Learning

```

1: Input: set of tasks  $\mathcal{V} = \{v_1, v_2, \dots\}$ , common information
2: Output: refined information  $\{\tilde{I}_v\}_{v \in \mathcal{V}}$ 
3: Seed selection: choose the root task  $s \in \mathcal{V}$ 
4: Tree construction: construct a rooted tree  $\mathcal{T} = (\mathcal{V}, \mathcal{E})$  defining
   the information flow
5: Budget allocation: assign refinement budgets  $\{b_v\}_{v \in \mathcal{V}}$  such
   that  $\sum_v b_v = B$ 
6: Cascaded transfer: traverse the tree in a topological (root-to-
   leaf) order; for each task  $v$ :
7:   If  $v = s$ , initialize  $I_v^{(0)}$  arbitrarily
8:   If  $v \neq s$ , let  $u = \text{pa}(v)$  and set  $I_v^{(0)} = \tilde{I}_u$ 
9:   Refine:  $\tilde{I}_v = G_v^{b_v}(I_v^{(0)})$ 
10: Return: refined information  $\{\tilde{I}_v : v \in \mathcal{V}\}$ 

```

in which tasks are trained sequentially following the tree structure: the root task is trained first, and each other task is initialized from its parent and then refined using a local optimization procedure.

Figure 1 provides an overview of the CTL framework, illustrating how, starting from a chosen seed, information is propagated from task to task along the cascade. The CTL framework involves three central design choices:

- (i) *Seed selection.* The root task  $s$  that initiates the cascade.
- (ii) *Graph construction.* The rooted tree defining the information flow between tasks.
- (iii) *Budget allocation and refinement.* The assignment of refinement budgets  $\{b_v\}_{v \in \mathcal{V}}$  and the execution of local optimization along the cascade.

Algorithm 1 summarizes the CTL process.

### 2.2. Cascade Construction Strategies

Cascaded Transfer Learning relies on the idea that tasks can benefit from information propagated through a predefined structure connecting them. Rather than treating transfers separately, CTL schedules task learning so that each task is refined using information inherited from neighboring tasks, under a global budget constraint. The central question is therefore not only how to define neighborhoods between tasks, but also under which conditions routing information through an intermediate task is beneficial.

At the local level, consider three tasks  $(S, I, T)$ , where  $S$  and  $T$  are neighbors according to a given task distance. A direct transfer  $S \rightarrow T$  is not necessarily optimal under a fixed training budget: in certain regimes, refining an intermediate task  $I$  and transferring along the path  $S \rightarrow I \rightarrow T$  can lead to improved performance. Section 3 characterizes sufficient conditions under which the error introduced by intermediate optimization remains controlled and can accumulate favorably along such paths. In particular, when local optimization procedures are stable and task distances are sufficiently small along the path, information can be propagated through cascades without degrading performance. When this behavior occurs systematically across a collection of tasks, it becomes natural to schedule transfers through a global structure that connects neighboring tasks while supporting propagation through intermediates. Rooted trees provide a simple and effective way to implement this idea: each task receives information from a limited set of upstream neighbors and contributes refined information downstream, ensuring a coherent and budget-aware diffusion process.

The specific choice of tree construction depends on how task distances are estimated and on the reliability of the available information. In practical settings, task relationships are typically inferred from noisy proxies such as empirical

165 features, gradients, or learned representations. In this context, distance-based constructions provide a robust way to  
 166 organize transfer. In particular, MSTs connect neighboring  
 167 tasks using pairwise distances while minimizing total edge  
 168 cost, encouraging locality and limiting the accumulation of  
 169 transfer error. As a result, MST-based cascades constitute a  
 170 strong default choice for CTL.  
 171

172 **Extensions.** The abstract CTL framework can be naturally  
 173 generalized in a number of directions that we mention here,  
 174 although their proper investigation is out of the scope of  
 175 this paper. First, while our presentation is restricted to a  
 176 single rooted tree, the same principles apply to a *cascade*  
 177 *forest* in which multiple seeds initiate independent cascades  
 178 over disjoint subsets of tasks. Deciding when more than  
 179 one seed would be beneficial and selecting those seeds,  
 180 is itself a design problem that can be informed by task  
 181 distances, clustering criteria, or transferability estimates.  
 182 Second, CTL does not rely on supervised losses: any setting  
 183 in which tasks share a common information space  $\mathcal{I}$  and  
 184 a local adaptation operator, can support cascaded transfer,  
 185 including unsupervised, self-supervised, or representation  
 186 learning scenarios. Third, Definition 2.1 can be stated more  
 187 generally for DAGs, which include rooted trees. That would  
 188 allow a task to fuse information from multiple parents.  
 189

### 3. Theoretical Analysis of Cascaded Transfer Learning

190 We justify CTL over rooted trees, regardless how the tree  
 191 is constructed. Our analysis shows that (i) tree-structured  
 192 cascades can improve over independent or pairwise transfer  
 193 by routing long-range transfers through sequences of shorter  
 194 ones in terms of task distance, (ii) only mild structural con-  
 195 ditions on the tree are required, and (iii) distance-based trees  
 196 such as MSTs naturally satisfy these conditions when task  
 197 distances are only imperfectly observed.  
 198

#### 3.1. Parameter-Space Analysis

199 For each task  $v \in \mathcal{V}$ , we associate a loss  $\mathcal{L}_v : \mathbb{R}^d \rightarrow \mathbb{R}$   
 200 and minimizer  $\theta_v^* \in \mathbb{R}^d$ . Let  $\eta \in (0, 1)$  be the learning  
 201 step associated to the refinement operator  $G_v$ . We assume  
 202 a contraction property in parameter space: for all  $v \in \mathcal{V}$ ,  
 203 there exists  $\rho_v \in (0, 1)$  such that for all  $(\theta, b) \in \mathbb{R}^d \times \mathbb{N}^*$ ,  
 204

$$\|G_v^b(\theta) - \theta_v^*\| \leq \rho_v^b \|\theta - \theta_v^*\|.$$

205 This holds, for instance, for gradient descent on strongly  
 206 convex and smooth losses. We assume a latent task geometry  
 207  $d(u, v) = \|\theta_u^* - \theta_v^*\|$ , as a proxy of the transfer difficulty  
 208 between tasks.  
 209

210 Let  $\mathcal{T}$  be a rooted tree with root  $s$ . Given a budget allo-  
 211 cation  $\{b_v\}_v$  with  $\sum_v b_v = B \in \mathbb{N}^*$ , CTL refines tasks in a  
 212

topological order:

$$\tilde{\theta}_s = G_s^{b_s}(\theta_{\text{init}}), \quad \tilde{\theta}_v = G_v^{b_v}(\tilde{\theta}_{\text{pa}(v)}), \quad v \neq s.$$

**Proposition 3.1** (Edge-wise propagation). *For any non-root node  $v$  with  $u = \text{pa}(v)$ , and for  $j, m \in \mathbb{N}^*$  such that  $S_{i,m} = \sum_{\ell=i}^m b_{v_\ell}$ ,*

$$\|\tilde{\theta}_v - \theta_v^*\| \leq \rho_v^{b_v} (\|\tilde{\theta}_u - \theta_u^*\| + d(u, v)).$$

*The recursion along the unique root-to- $v$  path yields*

$$\begin{aligned} \|\tilde{\theta}_v - \theta_v^*\| &\leq \prod_{i=1}^m \rho_{v_i}^{b_{v_i}} \|\tilde{\theta}_s - \theta_s^*\| \\ &+ \sum_{i=1}^m \rho_{v_i}^{S_{i,m}} d(v_{i-1}, v_i), \end{aligned}$$

*where each edge contribution is damped by downstream refinement.*

This bound shows that CTL improves over direct transfer when: (i) edges connect nearby tasks (locality), (ii) the structure is acyclic and ordered, (iii) budgets prevent error accumulation along deep paths. These conditions depend on edge lengths, not on how the tree is constructed.

**MST as a default.** When only pairwise dissimilarities are available, minimizing cumulative edge costs is a natural objective. Minimum spanning trees (MST) satisfy locality, are acyclic by construction, and are stable to noise in similarity estimates, making them a robust default choice.

**CTL vs TL. Why Cascades?** The central question is whether routing transfer through intermediate tasks can reduce the error induced by task mismatch, under a fixed refinement budget. The following result formalizes this intuition by isolating the *transfer bias* induced by task geometry.

**Theorem 3.2.** *Assume the seed task is solved exactly,  $\tilde{\theta}_s = \theta_s^*$ . Let  $(v_0 = s \rightarrow v_1 \rightarrow \dots \rightarrow v_m = v)$  be a path in a CTL tree, with uniform refinement budget  $b$  on each non-root node,  $\rho_{\max} = \max_{1 \leq i \leq m} \rho_{v_i}$ ,  $\delta_i = d(v_{i-1}, v_i)$ , and  $\delta_{\max} = \max_{1 \leq i \leq m} \delta_i$ .*

*Then CTL satisfies*

$$\|\tilde{\theta}_v^{\text{CTL}} - \theta_v^*\| \leq \sum_{i=1}^m \rho_{v_i}^{(m-i+1)b} \delta_i,$$

*whereas a direct transfer strategy satisfies*

$$\|\tilde{\theta}_v^{\text{TL}} - \theta_v^*\| \leq \rho_v^b d(s, v).$$

*Consequently, CTL yields a strictly smaller upper bound whenever*

$$\sum_{i=1}^m \rho_{v_i}^{(m-i+1)b} \delta_i < \rho_v^b d(s, v).$$

220 A sufficient condition is

$$221 \quad \delta_{\max}(1 - \rho_{\max}^{mb}) < d(s, v)(1 - \rho_{\max}^b).$$

222 Theorem 3.2 highlights a fundamental distinction between  
 223 cascaded and direct transfer. Star transfer incurs a *single*  
 224 geometric error proportional to the global distance  $d(s, v)$   
 225 between the seed and the target, which is discounted only  
 226 once by the local refinement. In contrast, CTL decomposes  
 227 this global mismatch into a sequence of local discrepancies  
 228  $\delta_i = d(v_{i-1}, v_i)$ , each of which is further attenuated by  
 229 all downstream refinements. As a result, when tasks vary  
 230 smoothly along the cascade (small  $\delta_i$ ), the discounted sum  
 231 of local errors is strictly smaller than the discounted global  
 232 jump. This explains why locality-preserving trees are effec-  
 233 tive: CTL exploits smoothness of the task geometry rather  
 234 than relying on a single global alignment.

### 235 3.2. Feature-Space Analysis

236 We now consider the linear feature-space setting with ran-  
 237 dom design and  $n \geq d$ . For each task  $v \in \mathcal{V}$ , let  $\mathbf{X}_v$  be  
 238 a random variable taking values in  $\mathbb{R}^{n \times d}$ , and let  $\mathbf{y}_v$  be  
 239 a random variable taking values in  $\mathbb{R}^n$ . We consider the  
 240 quadratic objective

$$241 \quad \mathcal{L}_v(\boldsymbol{\theta}) = \frac{1}{2} \|\mathbf{X}_v \boldsymbol{\theta} - \mathbf{y}_v\|^2, \quad \boldsymbol{\theta} \in \mathbb{R}^d.$$

242 A single gradient descent step with step size  $\eta > 0$  yields  
 243 the affine update  $G_v(\boldsymbol{\theta}) = \mathbf{M}_v \boldsymbol{\theta} + \eta \mathbf{X}_v^\top \mathbf{y}_v$ , where  $\mathbf{M}_v =$   
 244  $\mathbf{I}_d - \eta \mathbf{X}_v^\top \mathbf{X}_v$ .

245 We assume that  $\mathbf{X}_v^\top \mathbf{X}_v \succ 0$  almost surely and choose a  
 246 step size  $\eta \in (0, 2/\lambda_{\max}(\mathbf{X}_v^\top \mathbf{X}_v))$ , where  $\lambda_{\max}$  denotes  
 247 the largest eigenvalue. Under this condition, the linear operator  
 248  $\mathbf{M}_v = \mathbf{I}_d - \eta \mathbf{X}_v^\top \mathbf{X}_v$  satisfies  $\|\mathbf{M}_v\|_2 = \rho_v < 1$ ,  
 249 and therefore the refinement operator  $G_v$  is a contraction on  
 250  $(\mathbb{R}^d, \|\cdot\|_2)$  almost surely. We further assume a realizable  
 251 random-design linear model: for each task  $v \in \mathcal{V}$ , there  
 252 exists a parameter vector  $\boldsymbol{\theta}_v^* \in \mathbb{R}^d$  such that  $\mathbf{y}_v = \mathbf{X}_v \boldsymbol{\theta}_v^*$ .  
 253 In this setting, gradient-based refinement contracts toward  
 254 the task-specific optimum  $\boldsymbol{\theta}_v^*$ , and the same edge-wise and  
 255 path-wise propagation bounds as in parameter space apply  
 256 directly.

257 **Proposition 3.3** (Edge-wise propagation in feature space).  
 258 *With the same notations as in Proposition 3.1,*

$$259 \quad \|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| \leq \rho_v^{b_v} (\|\tilde{\boldsymbol{\theta}}_u - \boldsymbol{\theta}_u^*\| + d(u, v)).$$

260 *Unrolling this recursion along the unique root-to-v path*  
 261 *yields*

$$262 \quad \begin{aligned} \|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| &\leq \prod_{i=1}^m \rho_{v_i}^{b_{v_i}} \|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\| \\ 263 &+ \sum_{i=1}^m \rho_{v_i}^{S_{i,m}} d(v_{i-1}, v_i), \end{aligned}$$

264 where each edge contribution is damped by downstream  
 265 refinement.

### 266 3.3. Noisy Feature-Space Analysis

267 We consider a random-design linear regression model with  
 268 additive observation noise. For each task  $v \in \mathcal{V}$ , let  $\mathbf{X}_v$  be  
 269 a random variable taking values in  $\mathbb{R}^{n \times d}$ , and let  $\mathbf{y}_v$  be  
 270 a random variable taking values in  $\mathbb{R}^n$ . We assume that  
 271  $\mathbf{y}_v$  admits the decomposition  $\mathbf{y}_v = \mathbf{X}_v \boldsymbol{\theta}_v^* + \boldsymbol{\varepsilon}_v$ , where  
 272  $\boldsymbol{\theta}_v^* \in \mathbb{R}^d$  is a deterministic task-specific parameter and  
 273  $\boldsymbol{\varepsilon}_v$  is a noise random variable taking values in  $\mathbb{R}^n$  with  
 274 independent, mean-zero, sub-Gaussian coordinates.

275 We assume that the feature-space contraction condition  
 276 holds for a suitable step size  $\eta > 0$ . As in the previous  
 277 setting, gradient-based refinement induces a contraction  
 278 toward a task-dependent optimum. In the presence of ob-  
 279 servation noise, however, this contraction occurs toward the  
 280 empirical minimizer  $\tilde{\boldsymbol{\theta}}_v = \arg \min_{\boldsymbol{\theta} \in \mathbb{R}^d} \frac{1}{2} \|\mathbf{X}_v \boldsymbol{\theta} - \mathbf{y}_v\|_2^2$ ,  
 281 rather than toward the population parameter  $\boldsymbol{\theta}_v^*$ . The dis-  
 282 crepancy  $\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*$  captures statistical estimation error and  
 283 perturbs the task geometry induced by the feature space. In  
 284 expectation, the magnitude of this perturbation is controlled  
 285 by the conditioning of the task-specific design through the  
 286 linear operator  $\mathbf{A}_v = (\mathbf{X}_v^\top \mathbf{X}_v)^{-1} \mathbf{X}_v^\top$ , which governs how  
 287 observation noise propagates from data space to parameter  
 288 space.

289 **Proposition 3.4** (Expected noisy propagation along a CTL  
 290 tree). *Under the noisy linear model and the feature-space  
 291 contraction assumption, the expected error propagation  
 292 along a CTL tree satisfies the following recursion. With the  
 293 same notations as in Proposition 3.1,*

$$294 \quad \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|] \leq \rho_v^{b_v} (\mathbb{E}[\|\tilde{\boldsymbol{\theta}}_u - \boldsymbol{\theta}_u^*\|] + d(u, v)) \\ 295 \quad + \sigma_v (1 + \rho_v^{b_v}) \|\mathbf{A}_v\|_F.$$

296 *Unrolling this recursion along the root-to-v path ( $v_0 = s \rightarrow$   
 297  $\dots \rightarrow v_m = v$ ) yields*

$$298 \quad \begin{aligned} \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|] &\leq \left( \prod_{j=1}^m \rho_{v_j}^{b_{v_j}} \right) \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\|] \\ 299 &+ \sum_{i=1}^m \left( \prod_{j=i}^m \rho_{v_j}^{b_{v_j}} \right) d(v_{i-1}, v_i) \\ 300 &+ \sum_{i=1}^m \left( \prod_{j=i+1}^m \rho_{v_j}^{b_{v_j}} \right) \sigma_{v_i} (1 + \rho_{v_i}^{b_{v_i}}) \|\mathbf{A}_{v_i}\|_F. \end{aligned}$$

301 This bound shows that, in expectation, CTL decomposes the  
 302 final error into two discounted components: a bias term, and  
 303 a noise term capturing the propagation of estimation error.  
 304 As in the noiseless setting, locality reduces transfer bias,  
 305 while sufficient downstream refinement or shallow cascades  
 306 limit noise accumulation.

## 275 4. Cascade Construction Algorithm

276 Building on the theoretical analysis of Section 3, we de-  
 277 scribe a simple and scalable procedure for constructing tree-  
 278 structured cascades that organize information transfer under  
 279 a global budget constraint. The cascade is defined by select-  
 280 ing a representative seed task and a rooted tree that specifies  
 281 how transfer is routed across tasks.  
 282

### 283 4.1. Seed Selection and Tree Construction

284 The seed task  $s$  serves as the root of the cascade and in-  
 285 fluences all downstream transfers. To ensure robustness to  
 286 heterogeneity and noise, we select the *medoid* of the task  
 287 set, defined as the task minimizing the sum of distances to  
 288 all others. Unlike centroid-based choices, the medoid is  
 289 an actual task and provides a stable, parameter-free, and  
 290 outlier-resistant initialization across settings.  
 291

292 Given a notion of task distances, the cascade structure is  
 293 specified by a rooted tree that determines how information  
 294 propagates from the seed to the remaining tasks. We focus  
 295 on distance-based constructions, where a symmetric task  
 296 distance matrix is used to connect tasks through local neigh-  
 297 borhoods. In this setting, MSTs provide a robust default,  
 298 as they favor short transfers, limit error accumulation, and  
 299 remain stable under noisy distance estimates. The resulting  
 300 tree is rooted at the seed and oriented away from it to define  
 301 the cascade, while alternative construction strategies can be  
 302 substituted within the same CTL framework when richer  
 303 task information is available.  
 304

### 305 4.2. Computational Complexity

306 The computational cost of CTL is dominated by the con-  
 307 struction of the cascade structure. For distance-based  
 308 cascades such as MSTs, building a rooted tree over  $|\mathcal{V}|$   
 309 tasks from pairwise distances requires  $\mathcal{O}(|\mathcal{V}|^2 \log |\mathcal{V}|)$  time.  
 310 While quadratic in the number of tasks, this cost remains  
 311 tractable for the task collections considered in our ex-  
 312 periments and is incurred only once, prior to adaptation. Once  
 313 the cascade is constructed, training proceeds locally along  
 314 the tree: each task is refined exactly once from its parent,  
 315 and the total optimization cost scales linearly with the global  
 316 refinement budget. As a result, the dominant cost during  
 317 deployment is governed by the budgeted local updates rather  
 318 than by joint or repeated task optimization.  
 319

## 320 5. Experiments

321 The goal of our experiments is twofold: (i) to demonstrate  
 322 that CTL effectively exploits task-graph structure to improve  
 323 performance under a fixed or reduced computational budget,  
 324 and (ii) to show that these gains are not specific to a partic-  
 325 ular data modality or learning problem. We therefore evaluate  
 326

327 CTL across a diverse set of settings, including synthetic and  
 328 real-world regression, as well as image-based classification  
 329 tasks. Across all experiments, learning follows the same  
 330 paradigm. Each task is associated with a local objective and  
 331 is optimized independently using gradient-based methods.  
 332 Task interactions occur exclusively through parameter ini-  
 333 tialization: each node in the cascade is initialized from its  
 334 parent in the tree, after which local refinement is performed.  
 335 The task graph is constructed using training data only, while  
 336 all reported results are computed on held-out test data, en-  
 337 suring that task relationships are inferred without access to  
 338 evaluation samples.

### 339 5.1. Datasets

340 **Synthetic Dataset.** We construct families of linear regres-  
 341 sion tasks of the form  $\mathbf{y}_v = \mathbf{X}_v \boldsymbol{\theta}_v + \boldsymbol{\varepsilon}_v$ , where  $\mathbf{X}_v \in \mathbb{R}^{n \times d}$   
 342 has i.i.d. standard Gaussian rows and  $\boldsymbol{\varepsilon}_v \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}_n)$ . To  
 343 control task relatedness, we consider the following gener-  
 344 ative model: we first draw a global center  $\boldsymbol{\theta}_0$  and  $K$   
 345 cluster shifts  $\delta_c \sim \mathcal{N}(\mathbf{0}, \tau_{\text{between}}^2 \mathbf{I}_d)$ , every task  $v$  is as-  
 346 signed a cluster  $c(v)$  and its parameter is generated as  
 347  $\boldsymbol{\theta}_v = \boldsymbol{\theta}_0 + \delta_{c(v)} + \boldsymbol{\zeta}_v$  with  $\boldsymbol{\zeta}_v \sim \mathcal{N}(\mathbf{0}, \tau_{\text{within}}^2 \mathbf{I}_d)$ . Varying  $\tau$   
 348 or  $(\tau_{\text{within}}, \tau_{\text{between}})$  controls the degree and structure of het-  
 349 erogeneity, enabling systematic evaluation across regimes  
 350 ranging from nearly homogeneous tasks to well-separated  
 351 clusters, see Figure 3.

352 **Real Datasets.** We evaluate CTL on both time-series fore-  
 353 casting and image classification. For time-series regression,  
 354 we use the UK electricity dataset<sup>1</sup>, which provides aggregated  
 355 half-hourly residential demand from smart meters operated  
 356 by two major UK distribution network operators. We focus on a  
 357 subset of 100 nodes in the Oxford urban area and predict feeder-level  
 358 demand using calendar features and lagged load values. To assess  
 359 generality beyond regression, we also consider high-dimensional image  
 360 classification tasks derived from Fashion-MNIST and CIFAR-10.  
 361 Each task corresponds to a binary classification problem  
 362 defined over a pair of classes, and multiple related tasks are  
 363 obtained by training on different subsets of the data. This  
 364 setting yields collections of closely related but non-identical  
 365 tasks, making it well suited for evaluating cascaded transfer  
 366 learning in classification regimes.

### 367 5.2. Experimental Settings

368 For all regression experiments, datasets are split into dis-  
 369 joint training and test sets. In the synthetic setting, we use  
 370 fixed sizes of  $n_{\text{train}} = 64$  and  $n_{\text{test}} = 128$  per task. For the  
 371 real-world time-series dataset, the training set spans from  
 372 February 14, 2024, at 12:00 a.m. to February 24, 2024, at  
 373 11:30 p.m., while the test set spans from February 25, 2024,

<sup>1</sup><https://weave.energy/>

at 12:00 a.m. to February 28, 2024, at 11:30 p.m. For image classification experiments, each task is trained and evaluated using fixed-size train and test splits, and all methods use the same model architecture and optimization hyperparameters across tasks. In all cases, task graphs are constructed exclusively from training data. Models are evaluated on the test set using the same metrics (root mean squared error) across all methods, and all results are averaged over 50 random seeds to account for variability due to data sampling and initialization.

### 5.3. Baselines

We compare our approach with several tree-structured or non-transfer baselines. All methods use identical training routines, both for seed training and local refinement.

– *Individual models.* For each task, an individual model is trained independently for a fixed computational budget. This serves as a baseline that uses no multi-task or transfer learning.

– *Star Tree.* A seed task  $s$  is selected, and then a star graph is constructed:  $(s \rightarrow v)$ , for all  $v$ . Each node receives the seed parameters and performs independent refinement.

– *Random Tree.* A seed task  $s$  is selected, and a random spanning tree over all tasks is generated by sampling a Prüfer sequence uniformly (Kumar et al., 1998; Deo & Micikevicius, 2001), which yields a labeled tree drawn uniformly from the set of all spanning trees. The tree is then rooted at  $s$  by orienting all edges away from it. This baseline represents an uninformed hierarchical transfer structure that does not exploit any task-related information.

– *Minimum-Spanning Tree.* We consider several MST-based cascades that differ only in the task distance used to construct the tree. We evaluate distances based on target similarity, feature representations, optimization geometry (gradient or model-based), and distributional or representation-level discrepancies (e.g. KL, Wasserstein, MMD, CKA). Full definitions of the distance metrics are provided in the appendix.

We do not compare CTL to multi-task learning or deep neural network baselines, as these methods optimize a fundamentally different objective through repeated joint training and shared representations, typically assuming abundant compute and task revisiting. In contrast, CTL targets frugal many-task adaptation, where each task is refined once under a strict global budget. To preserve alignment between theory and experiments, we therefore focus on linear and convex models, where contraction and error propagation can be characterized explicitly and resource usage is controlled.

### 5.4. Results

Table 2 and Figure 2 provide a unified view of all experiments across synthetic regression, UK electricity forecast-

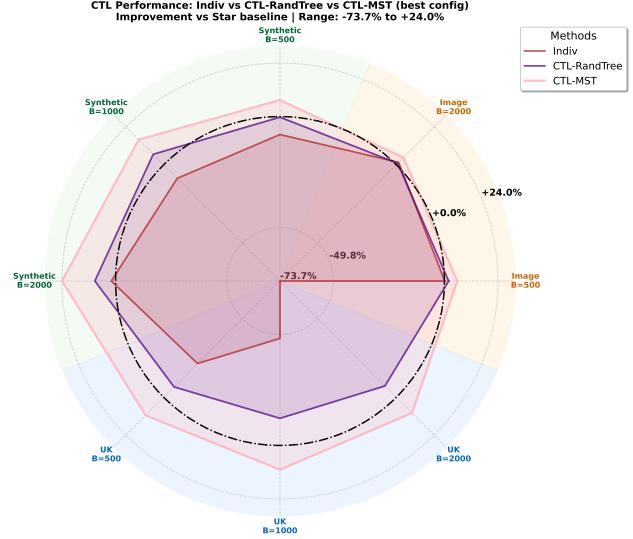


Figure 2. Radar plot of average percentage improvement over direct transfer training across datasets and budgets.

ing, and image classification, enabling direct comparison of transfer structures, task geometries, and budgets.

**Synthetic regression.** Across all regimes, cascaded transfer consistently outperforms independent training and shallow baselines, confirming that exploiting task structure remains beneficial even under increasing heterogeneity. At low to moderate variability, distances aligned with the underlying optimization or model geometry (for instance, gradient and parameter-based) yield the lowest errors across budgets, leading to improvements of 20–35% relative to independent training. As heterogeneity increases ( $\tau = 10$ ), absolute gains diminish and variance grows, reflecting weaker task relatedness; nevertheless, optimization-aligned distances remain systematically superior to purely feature-based or uninformed measures. Distributional distances provide competitive but less stable performance, while feature-based distances degrade more rapidly with increasing dispersion.

**UK electricity forecasting.** Across all settings, independent training yields the highest errors, highlighting the limitations of learning tasks in isolation. Shallow star-shaped transfer improves performance but exhibits higher variance and degrades in low-budget regimes, indicating sensitivity to unfavorable source–target mismatches. In contrast, cascaded transfer consistently reduces RMSE across budgets, with distances aligned to target or optimization geometry delivering the most stable gains. At low budget, cascaded transfer reduces RMSE by more than 25% relative to independent training, while at higher budget reductions approach 50%. The aggregate results in Figure 2 confirm that geometry-aware distances consistently outperform unstructured or feature-only measures, demonstrating that cascaded transfer effectively amortizes optimization effort across re-

385  
 386 *Table 1.* Results with MST-based cascades aggregated by distance family (Feature, Target, Optimization). Values are averaged over  
 387 distance variants within each family. Lower is better for RMSE and higher is better for accuracy. **Bold** and underlined entries denote best  
 388 and second-best performance *within each row*.

Dataset	Params		Indiv	Baselines		CTL-RandTree	Feature	CTL-MST	
	T	B		Star	CTL-RandTree			Target	Optim
<b>Synthetic Regression (RMSE)</b>									
Syn-10	200	500	913.8 ± 168.4	845.5 ± 178.8	872.5 ± 253.1	840 ± 215	853 ± 228	<b>825 ± 221</b>	
Syn-10	200	1000	871.5 ± 141.6	802.4 ± 189.5	749.2 ± 171.0	752 ± 188	758 ± 176	<u>694 ± 161</u>	
Syn-10	200	2000	912.8 ± 174.3	930.2 ± 262.4	887.2 ± 246.1	846 ± 228	809 ± 221	<b>715 ± 178</b>	
<b>UK Electricity Forecasting (RMSE, kWh)</b>									
UK	100	500	2729.4 ± 8.8	2247.9 ± 51.8	2395.9 ± 59.4	2391 ± 51	2086 ± 49	<b>2040 ± 42</b>	
UK	100	1000	2694.4 ± 9.5	1820.7 ± 38.7	2042.4 ± 36.4	1993 ± 38	<u>1695 ± 26</u>	<b>1675 ± 26</b>	
UK	100	2000	2612.8 ± 7.3	1504.1 ± 19.0	1612.4 ± 16.3	1612 ± 19	<u>1388 ± 8</u>	<b>1375 ± 7</b>	
<b>Image Classification (Accuracy, %)</b>									
FMNIST	200	500	90.6 ± 5.2	90.5 ± 5.2	91.5 ± 4.1	<b>93.6 ± 3.6</b>	90.4 ± 4.9	92.8 ± 4.2	
FMNIST	200	2000	94.5 ± 3.0	94.4 ± 3.3	94.1 ± 3.2	<b>95.1 ± 3.1</b>	94.0 ± 3.7	<u>95.1 ± 3.0</u>	
CIFAR	200	500	60.9 ± 4.7	64.0 ± 4.3	65.2 ± 3.9	<u>66.0 ± 3.8</u>	64.8 ± 3.8	<b>66.6 ± 3.7</b>	
CIFAR	200	2000	67.7 ± 3.9	66.7 ± 3.8	67.3 ± 3.5	<u>67.9 ± 3.6</u>	67.1 ± 3.9	<b>68.6 ± 3.7</b>	

403 lated forecasting tasks under tight computational constraints.

404  
 405  
 406  
 407 **Image classification.** On Fashion-MNIST, cascaded transfer  
 408 consistently improves over independent training and  
 409 shallow baselines across all budgets. Distances based on  
 410 gradients or task-level feature statistics yield the most stable  
 411 improvements, particularly in low-budget regimes where  
 412 cascaded initialization compensates for limited task-specific  
 413 optimization. In contrast, random hierarchies and star-  
 414 shaped transfer provide only marginal gains, underscoring  
 415 the importance of task geometry beyond simple informa-  
 416 tion sharing. On CIFAR-10, improvements are smaller and  
 417 variability is higher, reflecting increased task difficulty and  
 418 weaker inter-task alignment. Nevertheless, cascaded transfer  
 419 remains beneficial: geometry-aware distances consistently  
 420 outperform independent training and unstructured baselines  
 421 across most settings. While no single distance dominates  
 422 uniformly, the results confirm that exploiting task geom-  
 423 etry through cascaded transfer yields systematic accuracy  
 424 gains even in challenging, high-dimensional classification  
 425 regimes.

426  
 427 **Summary.** Across regression, forecasting, and classifi-  
 428 cation tasks, the results confirm that organizing transfer  
 429 through a tree yields robust and consistent improvements  
 430 over both no-transfer and shallow-transfer baselines. Within  
 431 cascaded transfer, the choice of task distance is critical: dis-  
 432 tances aligned with the underlying optimization or target  
 433 geometry consistently outperform purely feature-based or  
 434 uninformed measures across datasets and budgets. These  
 435 findings closely match the theoretical analysis, which pre-  
 436 dicted that routing transfer through local, geometry-aware  
 437 cascades mitigates the bias incurred by long-range transfer  
 438 under constrained budgets.

## 6. Conclusion and Perspectives

We introduced Cascaded Transfer Learning, a framework for many-task learning under a strict global budget. CTL formulates knowledge transfer as a structured propagation process over a rooted tree, in which tasks are refined sequentially rather than jointly optimized. This design enables controlled information flow across tasks and departs from both independent training and shallow transfer schemes.

We provided a theoretical analysis of error propagation along cascades, establishing sufficient conditions under which routing transfer through intermediate tasks improves over direct transfer. The analysis clarifies how task distances, local optimization stability, and cascade depth jointly govern transfer bias and variance, and applies broadly across parameter and feature-space settings, including noisy regimes. Empirically, CTL consistently outperforms independent and star-shaped baselines on synthetic regression, electricity forecasting, and image classification. Simple geometry-based constructions such as minimum spanning trees emerge as robust and effective defaults, yielding reliable gains under realistic heterogeneity without additional modeling complexity.

Future work includes extending CTL to nonlinear models and deep representations, as well as to more general cascade structures such as directed acyclic graphs. Developing adaptive strategies for allocating refinement budgets based on task difficulty or uncertainty also represents a promising direction. Overall, CTL offers a principled and computationally efficient approach to structured knowledge transfer, and provides a foundation for further research on resource-aware many-task learning.

440  
441  
442  
443  
444  
445  
446  
447  
448  
**Impact Statement**

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here.

449  
450  
451  
452  
453  
**References**

Aribandi, V., Tay, Y., Schuster, T., Rao, J., Zheng, H. S., Mehta, S. V., Zhuang, H., Tran, V. Q., Bahri, D., Ni, J., et al. ExT5: Towards extreme multi-task scaling for transfer learning. In *International Conference on Learning Representations*, 2022.

Deo, N. and Micikevicius, P. Prufer-like codes for labeled trees. *Congressus Numerantium*, pp. 65–74, 2001.

Hashimoto, K., Xiong, C., Tsuruoka, Y., and Socher, R. A joint many-task model: Growing a neural network for multiple NLP tasks. In *Conference on Empirical Methods in Natural Language Processing*, pp. 1923–1933, 2017.

He, X., Alesiani, F., and Shaker, A. Efficient and scalable multi-task regression on massive number of tasks. In *AAAI Conference on Artificial Intelligence*, volume 33, pp. 3763–3770, 2019.

Kumar, V., Deo, N., and Kumar, N. Parallel generation of random trees and connected graphs. *Congressus Numerantium*, pp. 7–18, 1998.

Liu, J., Xia, Z., Lei, Y., Li, X., and Wang, X. Multi-faceted hierarchical multi-task learning for a large number of tasks with multi-dimensional relations. *Preprint arXiv:2110.13365*, 2021.

Nie, Y., Nguyen, N. H., Sinthong, P., and Kalagnanam, J. A time series is worth 64 words: Long-term forecasting with transformers, 2022.

Pan, S. J. and Yang, Q. A survey on transfer learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10):1345–1359, 2009.

Ruiz, C., Alaíz, C. M., and Dorronsoro, J. R. A survey on kernel-based multi-task learning. *Neurocomputing*, 577: 127–255, 2024.

Strezoski, G., Noord, N. v., and Worring, M. Many task learning with task routing. In *IEEE/CVF International Conference on Computer Vision*, pp. 1375–1384, 2019.

Tan, Y., Zhang, E., Li, Y., Huang, S.-L., and Zhang, X.-P. Transferability-guided cross-domain cross-task transfer learning. *IEEE Transactions on Neural Networks and Learning Systems*, 2024.

Weiss, K., Khoshgoftaar, T. M., and Wang, D. A survey of transfer learning. *Journal of Big data*, 3(1):9, 2016.

Woo, G., Liu, C., Kumar, A., Xiong, C., Savarese, S., and Sahoo, D. Unified training of universal time series forecasting transformers, 2024.

Zhang, Y. and Yang, Q. An overview of multi-task learning. *National Science Review*, 5(1):30–43, 2018.

Zhang, Y. and Yang, Q. A survey on multi-task learning. *IEEE Transactions on Knowledge and Data Engineering*, 34(12):5586–5609, 2021.

Zhuang, F., Qi, Z., Duan, K., Xi, D., Zhu, Y., Zhu, H., Xiong, H., and He, Q. A comprehensive survey on transfer learning. *Proceedings of the IEEE*, 109(1):43–76, 2020.

## A. Proofs for Theoretical Analysis of Cascaded Transfer Learning

### A.1. Parameter-Space

498 *Edge-wise propagation in parameter space.* By definition of the CTL update,  $\tilde{\boldsymbol{\theta}}_v = G_v^{b_v}(\tilde{\boldsymbol{\theta}}_{\text{pa}(v)})$ . By the contractive  
 499 refinement assumption,

$$500 \quad 501 \quad 502 \quad \|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| = \|G_v^{b_v}(\tilde{\boldsymbol{\theta}}_{\text{pa}(v)}) - \boldsymbol{\theta}_v^*\| \leq \rho_v^{b_v} \|\tilde{\boldsymbol{\theta}}_{\text{pa}(v)} - \boldsymbol{\theta}_v^*\|.$$

Applying the triangle inequality,

$$503 \quad 504 \quad \|\tilde{\boldsymbol{\theta}}_{\text{pa}(v)} - \boldsymbol{\theta}_v^*\| \leq \|\tilde{\boldsymbol{\theta}}_{\text{pa}(v)} - \boldsymbol{\theta}_{\text{pa}(v)}^*\| + d(\text{pa}(v), v).$$

505 Combining the two inequalities yields the result.  $\square$

507 *Path-wise error decomposition.* Applying the above to node  $v_1$ :

$$509 \quad 510 \quad \|\tilde{\boldsymbol{\theta}}_{v_1} - \boldsymbol{\theta}_{v_1}^*\| \leq \rho_{v_1}^{b_{v_1}} (\|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\| + d(s, v_1)).$$

511 Same with node  $v_2$  and substitute the previous inequality:

$$513 \quad \|\tilde{\boldsymbol{\theta}}_{v_2} - \boldsymbol{\theta}_{v_2}^*\| \leq \rho_{v_1}^{b_{v_1}} \rho_{v_2}^{b_{v_2}} \|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\| + \rho_{v_1}^{b_{v_1}} \rho_{v_2}^{b_{v_2}} d(s, v_1) + \rho_{v_2}^{b_{v_2}} d(v_1, v_2).$$

515 Proceeding inductively, each application introduces a multiplicative factor  $\rho_{v_i}^{b_{v_i}}$  on all upstream terms and adds a new edge  
 516 term. Hence, we get

$$518 \quad 519 \quad \|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| \leq \prod_{i=1}^m \rho_{v_i}^{b_{v_i}} \|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\| + \sum_{i=1}^m \rho_{v_i}^{S_{i,m}} d(v_{i-1}, v_i).$$

521 Setting  $\tilde{\boldsymbol{\theta}}_s = \boldsymbol{\theta}_s^*$  cancels out the bias term, enforcing  $b_{v_i} = b$  for all  $1 \leq i \leq m$ , plus bounding all  $\rho_{v_i}$  by  $\rho_{\max}$  and  $\delta_i$  by  
 522  $\delta_{\max}$  yields a finite geometric series of finite ratio  $\rho_{\max}^b$ .  $\square$

### A.2. Feature-Space

526 *Feature-space contraction.* The gradient update is  $G_v(\boldsymbol{\theta}) = \boldsymbol{\theta} - \eta \mathbf{X}_v^\top (\mathbf{X}_v \boldsymbol{\theta} - \mathbf{y}_v) = \mathbf{M}_v \boldsymbol{\theta} + \eta \mathbf{X}_v^\top \mathbf{y}_v$ . Since  $\mathbf{y}_v = \mathbf{X}_v \boldsymbol{\theta}_v^*$   
 527 and  $\boldsymbol{\theta}_v^*$  satisfies  $\mathbf{X}_v^\top \mathbf{X}_v \boldsymbol{\theta}_v^* = \mathbf{X}_v^\top \mathbf{y}_v$ , then  $G_v(\boldsymbol{\theta}_v^*) = \boldsymbol{\theta}_v^*$ , so  $\boldsymbol{\theta}_v^*$  is a fixed point. Thus,  $G_v(\boldsymbol{\theta}) - \boldsymbol{\theta}_v^* = \mathbf{M}_v(\boldsymbol{\theta} - \boldsymbol{\theta}_v^*)$ . Iterating yields the claim.  $\square$

530 *Edge-wise propagation in feature space.* The above yields  $\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^* = \mathbf{M}_v^{b_v}(\tilde{\boldsymbol{\theta}}_{\text{pa}(v)} - \boldsymbol{\theta}_v^*)$ . Taking norms and using  $\|\mathbf{M}_v^{b_v}\| \leq \rho_v^{b_v}$ ,

$$532 \quad \|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| \leq \rho_v^{b_v} \|\tilde{\boldsymbol{\theta}}_{\text{pa}(v)} - \boldsymbol{\theta}_v^*\|.$$

533 The result follows by the triangle inequality exactly as in the parameter-space case.  $\square$

535 *Path-wise error propagation in feature space.* The argument follows verbatim from the parameter-space analysis. Unrolling  
 536 the edge-wise inequality along the unique root-to- $v$  path yields the stated bound, with each edge contribution geometrically  
 537 damped by downstream refinements.  $\square$

### A.3. Noisy Feature-Space

541 *Empirical optimum.* The minimizer  $\hat{\boldsymbol{\theta}}_v$  of  $\frac{1}{2} \|\mathbf{X}_v \boldsymbol{\theta} - \mathbf{y}_v\|^2$  satisfies the equations

$$543 \quad \mathbf{X}_v^\top \mathbf{X}_v \hat{\boldsymbol{\theta}}_v = \mathbf{X}_v^\top \mathbf{y}_v = \mathbf{X}_v^\top (\mathbf{X}_v \boldsymbol{\theta}_v^* + \boldsymbol{\varepsilon}_v) = \mathbf{X}_v^\top \mathbf{X}_v \boldsymbol{\theta}_v^* + \mathbf{X}_v^\top \boldsymbol{\varepsilon}_v.$$

544 Since  $\mathbf{X}_v^\top \mathbf{X}_v \succ 0$ ,

$$545 \quad \hat{\boldsymbol{\theta}}_v = \boldsymbol{\theta}_v^* + (\mathbf{X}_v^\top \mathbf{X}_v)^{-1} \mathbf{X}_v^\top \boldsymbol{\varepsilon}_v.$$

547 Moreover, the gradient update is  $G_v(\boldsymbol{\theta}) = \mathbf{M}_v \boldsymbol{\theta} + \eta \mathbf{X}_v^\top \mathbf{y}_v$ , with  $\mathbf{M}_v = \mathbf{I}_d - \eta \mathbf{X}_v^\top \mathbf{X}_v$ . Using  $\mathbf{X}_v^\top \mathbf{y}_v = \mathbf{X}_v^\top \mathbf{X}_v \hat{\boldsymbol{\theta}}_v$ , we  
 548 have that  $\hat{\boldsymbol{\theta}}_v$  is a fixed point. Hence  $G_v(\boldsymbol{\theta}) - \hat{\boldsymbol{\theta}}_v = \mathbf{M}_v(\boldsymbol{\theta} - \hat{\boldsymbol{\theta}}_v)$ , and iterating gives  $G_v^{b_v}(\boldsymbol{\theta}) - \hat{\boldsymbol{\theta}}_v = \mathbf{M}_v^{b_v}(\boldsymbol{\theta} - \hat{\boldsymbol{\theta}}_v)$ .  $\square$

550 *Expected estimation error.* Let  $\mathbf{A}_v = (\mathbf{X}_v^\top \mathbf{X}_v)^{-1} \mathbf{X}_v^\top$ . From the empirical optimum lemma,  $\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^* = \mathbf{A}_v \boldsymbol{\epsilon}_v$ . Using  
 551 Jensen's inequality yields  $\mathbb{E}[\|\mathbf{A}_v \boldsymbol{\epsilon}_v\|] \leq \sqrt{\mathbb{E}[\|\mathbf{A}_v \boldsymbol{\epsilon}_v\|^2]}$ . Now,  
 552

$$553 \quad 554 \quad \|\mathbf{A}_v \boldsymbol{\epsilon}_v\|^2 = \boldsymbol{\epsilon}_v^\top \mathbf{A}_v^\top \mathbf{A}_v \boldsymbol{\epsilon}_v = \sum_{i,j} (\mathbf{A}_v^\top \mathbf{A}_v)_{ij} \boldsymbol{\epsilon}_{v,i} \boldsymbol{\epsilon}_{v,j}.$$

555 Since the coordinates are independent and mean-zero,  $\mathbb{E}[\boldsymbol{\epsilon}_{v,i} \boldsymbol{\epsilon}_{v,j}] = 0$  for  $i \neq j$ . Therefore,  
 556

$$557 \quad \mathbb{E}[\|\mathbf{A}_v \boldsymbol{\epsilon}_v\|^2] = \sum_i (\mathbf{A}_v^\top \mathbf{A}_v)_{ii} \mathbb{E}[\boldsymbol{\epsilon}_{v,i}^2] \leq \sigma_v^2 \sum_i (\mathbf{A}_v^\top \mathbf{A}_v)_{ii} = \sigma_v^2 \text{tr}(\mathbf{A}_v^\top \mathbf{A}_v) = \sigma_v^2 \|\mathbf{A}_v\|_F^2.$$

560 Combining yields  $\mathbb{E}[\|\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|] \leq \sigma_v \|\mathbf{A}_v\|_F$ . □  
 561

562 *Expected noisy edge-wise propagation.* With the contraction assumption for task  $v$ , we get  
 563

$$564 \quad 565 \quad \|\tilde{\boldsymbol{\theta}}_v - \hat{\boldsymbol{\theta}}_v\| \leq \rho_v^{b_v} \|\tilde{\boldsymbol{\theta}}_u - \hat{\boldsymbol{\theta}}_v\|.$$

566 Then, adding and subtracting  $\boldsymbol{\theta}_v^*$ ,  
 567

$$568 \quad \|\tilde{\boldsymbol{\theta}}_u - \hat{\boldsymbol{\theta}}_v\| \leq \|\tilde{\boldsymbol{\theta}}_u - \boldsymbol{\theta}_v^*\| + \|\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| \leq \|\tilde{\boldsymbol{\theta}}_u - \boldsymbol{\theta}_u^*\| + \|\boldsymbol{\theta}_u^* - \boldsymbol{\theta}_v^*\| + \|\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|.$$

569 Therefore,  
 570

$$571 \quad \|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| \leq \|\tilde{\boldsymbol{\theta}}_v - \hat{\boldsymbol{\theta}}_v\| + \|\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\| \leq \rho_v^{b_v} (\|\tilde{\boldsymbol{\theta}}_u - \boldsymbol{\theta}_u^*\| + \|\boldsymbol{\theta}_u^* - \boldsymbol{\theta}_v^*\| + \|\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|) + \|\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|.$$

573 Taking expectations and using  $\mathbb{E}[\|\hat{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|] \leq \sigma_v \|\mathbf{A}_v\|_F$  gives  
 574

$$575 \quad 576 \quad \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|] \leq \rho_v^{b_v} (\mathbb{E}[\|\tilde{\boldsymbol{\theta}}_u - \boldsymbol{\theta}_u^*\|] + \|\boldsymbol{\theta}_u^* - \boldsymbol{\theta}_v^*\|) + \sigma_v (1 + \rho_v^{b_v}) \|\mathbf{A}_v\|_F,$$

577 as claimed. □  
 578

579 *Expected path-wise propagation.* Let  $(v_0 = s \rightarrow v_1 \rightarrow \dots \rightarrow v_m = v)$  denote the unique root-to- $v$  path. By the expected  
 580 noisy edge-wise propagation lemma, for each  $i \geq 1$ ,  
 581

$$582 \quad 583 \quad \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_{v_i} - \boldsymbol{\theta}_{v_i}^*\|] \leq \rho_{v_i}^{b_{v_i}} (\mathbb{E}[\|\tilde{\boldsymbol{\theta}}_{v_{i-1}} - \boldsymbol{\theta}_{v_{i-1}}^*\|] + \|\boldsymbol{\theta}_{v_{i-1}}^* - \boldsymbol{\theta}_{v_i}^*\|) + \sigma_{v_i} (1 + \rho_{v_i}^{b_{v_i}}) \|\mathbf{A}_{v_i}\|_F.$$

584 Applying this inequality to  $v_1$  yields  
 585

$$586 \quad 587 \quad \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_{v_1} - \boldsymbol{\theta}_{v_1}^*\|] \leq \rho_{v_1}^{b_{v_1}} (\mathbb{E}[\|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\|] + \|\boldsymbol{\theta}_s^* - \boldsymbol{\theta}_{v_1}^*\|) + \sigma_{v_1} (1 + \rho_{v_1}^{b_{v_1}}) \|\mathbf{A}_{v_1}\|_F.$$

588 Applying the same inequality to  $v_2$  and bounding the expectation term,  
 589

$$590 \quad \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_{v_2} - \boldsymbol{\theta}_{v_2}^*\|] \leq \rho_{v_2}^{b_{v_2}} \rho_{v_1}^{b_{v_1}} \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\|] + \rho_{v_2}^{b_{v_2}} \rho_{v_1}^{b_{v_1}} \|\boldsymbol{\theta}_s^* - \boldsymbol{\theta}_{v_1}^*\| + \rho_{v_2}^{b_{v_2}} \|\boldsymbol{\theta}_{v_1}^* - \boldsymbol{\theta}_{v_2}^*\| + \rho_{v_2}^{b_{v_2}} \sigma_{v_1} (1 + \rho_{v_1}^{b_{v_1}}) \|\mathbf{A}_{v_1}\|_F \\ 591 \quad + \sigma_{v_2} (1 + \rho_{v_2}^{b_{v_2}}) \|\mathbf{A}_{v_2}\|_F.$$

593 Proceeding inductively, each step multiplies all upstream terms by  $\rho_{v_i}^{b_{v_i}}$  and adds a new distance term discounted by  
 594 downstream contractions, together with anew noise contribution discounted only by refinements performed after node  $v_i$ .  
 595

596 Collecting terms yields  
 597

$$598 \quad \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_v - \boldsymbol{\theta}_v^*\|] \leq \left( \prod_{j=1}^m \rho_{v_j}^{b_{v_j}} \right) \mathbb{E}[\|\tilde{\boldsymbol{\theta}}_s - \boldsymbol{\theta}_s^*\|] + \sum_{i=1}^m \left( \prod_{j=i}^m \rho_{v_j}^{b_{v_j}} \right) \|\boldsymbol{\theta}_{v_{i-1}}^* - \boldsymbol{\theta}_{v_i}^*\| \\ 599 \quad + \sum_{i=1}^m \left( \prod_{j=i+1}^m \rho_{v_j}^{b_{v_j}} \right) \sigma_{v_i} (1 + \rho_{v_i}^{b_{v_i}}) \|\mathbf{A}_{v_i}\|_F,$$

600 with the convention that empty products equal 1. This concludes the proof. □  
 601  
 602

---

605 **Algorithm 2** CTL Adaptation

---

606 1: **Input:** task observations  $\{\mathbf{y}_v\}_{v \in \mathcal{V}}$ , optional cluster count  $n_{clust}$ , clustering routine  $\text{findClusters}()$ , seed selection routine  
607  $\text{findSeed}()$ , seed-training budget  $b$ , per-cluster adaptation budget  $B$ , training operators  $\mathcal{A}_0$  (seed training) and  $\mathcal{A}_1$  (task adaptation).

608 2: **Output:** adapted models  $\{f_{\bar{\theta}_v}\}_{v \in \mathcal{V}}$

---

609 3: ■ (Optional) Cluster the tasks.

610 4: **if**  $K$  is specified **then**

611 5:    $\{C_k\}_{k=1}^{n_{clust}} \leftarrow \text{findClusters}(\{\mathbf{y}_v\}_{v \in \mathcal{V}}, n_{clust})$

612 6: **else**

613 7:    $K \leftarrow 1$ ;  $C_1 \leftarrow \mathcal{V}$

614 8: **end if**

615 9: ■ Process each cluster independently

616 10: **for each** cluster  $C$  **do**

617 11:    $s \leftarrow \text{findSeed}(C)$  ▷ Select cluster seed

618 12:    $(\mathcal{T}, E) \leftarrow \text{computeTree}(C, s)$

619 13: ■ Allocate adaptation budgets along the tree

620 14:    $\mathbf{w} \leftarrow \text{extractTreeWeights}(\mathcal{T}, E)$

621 15:    $\mathbf{w}' \leftarrow \text{normalize}(\mathbf{w})$

622 16:   **for**  $v \in C$  **do**

623 17:      $b_v \leftarrow \lceil w'_v \cdot B \rceil$

624 18:   **end for**

625 19: ■ Train the seed model

626 20:    $f_{\bar{\theta}_s} \leftarrow \mathcal{A}_0(f_{\theta}, \mathbf{y}_s, b)$

627 21: ■ Hierarchical model adaptation

628 22:   Initialize queue  $\mathcal{Q} \leftarrow [s]$

629 23:   **while**  $\neg \text{isEmpty}(\mathcal{Q})$  **do**

630 24:      $u \leftarrow \text{dequeue}(\mathcal{Q})$

631 25:     **for each** child  $v$  of  $u$  in  $(\mathcal{T}, E)$  **do** ▷ Adapt task  $v$  from its parent's model

632 26:        $\mathcal{Q} \leftarrow \text{enqueue}(\mathcal{Q}, v)$

633 27:        $f_{\bar{\theta}_v} \leftarrow \mathcal{A}_1(f_{\bar{\theta}_u}, \mathbf{y}_v, b_v)$

634 28:     **end for**

635 29:   **end while**

636 30: **end for**

637 31: **return**  $\{f_{\bar{\theta}_v}\}_{v \in \mathcal{V}}$ .

---

## B. Cascaded Adaptation Algorithm

638 Algorithm 2 implements the cascaded adaptation procedure used throughout the paper. The algorithm performs hierarchical  
639 task adaptation under a global optimization budget by propagating information along a directed tree structure. Tasks  
640 may optionally be partitioned into clusters to limit transfer across distant regimes and reduce the risk of negative transfer.  
641 Within each cluster, a tree-structured cascade is constructed, and adaptation proceeds from a designated seed task toward  
642 downstream tasks. The seed task is trained from scratch, while each remaining task is adapted exactly once from its parent  
643 in the tree.

644 Adaptation follows a top-down traversal of the tree, ensuring that parent tasks are fully adapted before their children. This  
645 enforces an acyclic flow of information and prevents backward or joint updates, which is consistent with the cascade analysis  
646 developed in the main text. The specific form of the local adaptation operator is left abstract, allowing instantiation with  
647 linear or nonlinear models, provided that adaptation remains task-local. A fixed global budget  $B$  is allocated across tasks  
648 using normalized tree-based weights. While the choice of weighting scheme is flexible, it reflects the intuition that tasks  
649 deeper in the cascade or further from the seed may require additional adaptation effort. In practice, discretization effects may  
650 introduce a small budget slack, which is negligible relative to  $B$  and does not affect the qualitative behavior of the algorithm.

651 Overall, this procedure provides a practical realization of cascaded transfer learning, emphasizing budget awareness, stability,  
652 and modularity in large collections of related tasks.

## C. Distance metrics for MST-based cascades

653 We group the task distance metrics used to construct MST-based cascades into the following families:  
654

Table 2. Unified experimental results across all datasets. **Synthetic regression:** Test RMSE. **UK electricity:** Test RMSE (kWh). **Image classification:** Accuracy (%). Each configuration is evaluated across 50 simulations. **Bold** and underlined entries denote best and second-best performance *within each row*.

Dataset	Params	$T$	$B$	Indiv	Baselines	CTL-RandTree	CTL-MST variants					JS	CKA		
							Feature	Target	Gradient	Model	KL	Wasserstein			
<b>Synthetic Regression (RMSE)</b>															
Syn-2	200	500	494.2 ± 149.9	440.4 ± 155.9	455.4 ± 192.2	447.9 ± 166.4	445.5 ± 169.1	432.4 ± 153.3	<b>407.7 ± 143.3</b>	453.7 ± 156.4	417.4 ± 143.0	460.8 ± 197.7	485.6 ± 254.9	458.4 ± 203.1	
Syn-2	200	1000	458.9 ± 134.2	382.0 ± 139.1	386.1 ± 146.5	363.2 ± 131.7	373.3 ± 121.2	368.8 ± 143.5	<u>356.9 ± 151.3</u>	369.5 ± 132.6	412.8 ± 150.7	400.5 ± 156.9	<b>355.6 ± 130.9</b>	382.7 ± 114.6	385.5 ± 171.9
Syn-2	200	2000	499.2 ± 178.4	451.6 ± 181.9	409.8 ± 162.5	448.8 ± 187.7	418.4 ± 152.5	370.1 ± 163.8	<b>347.6 ± 144.2</b>	421.5 ± 181.3	433.0 ± 178.4	428.3 ± 181.5	429.0 ± 191.7	457.6 ± 215.2	403.9 ± 167.0
Syn-5	200	500	595.9 ± 166.3	547.8 ± 224.6	549.4 ± 164.9	528.9 ± 209.7	<b>522.1 ± 174.2</b>	525.7 ± 176.0	532.0 ± 165.3	543.5 ± 198.8	530.4 ± 173.5	546.5 ± 181.0	537.9 ± 180.4	546.0 ± 171.0	547.8 ± 187.3
Syn-5	200	1000	550.6 ± 130.6	<b>484.7 ± 142.6</b>	460.1 ± 136.7	449.8 ± 149.6	450.7 ± 158.0	441.4 ± 150.5	444.5 ± 150.6	460.4 ± 154.9	472.3 ± 164.9	476.2 ± 151.7	466.9 ± 144.3	<b>437.3 ± 149.5</b>	472.9 ± 173.6
Syn-5	200	2000	595.0 ± 185.7	541.2 ± 187.1	520.6 ± 204.2	517.9 ± 182.0	507.8 ± 190.2	<b>437.4 ± 152.5</b>	447.1 ± 148.3	511.0 ± 176.4	517.7 ± 184.0	543.5 ± 213.1	517.3 ± 194.4	517.3 ± 184.0	545.2 ± 199.5
Syn-10	200	500	913.8 ± 168.4	845.5 ± 178.8	872.5 ± 253.1	835.9 ± 189.5	857.8 ± 215.5	844.7 ± 254.5	<b>805.2 ± 187.8</b>	891.9 ± 262.6	860.7 ± 257.7	823.8 ± 227.4	855.7 ± 226.6	<b>800.4 ± 175.3</b>	845.9 ± 213.5
Syn-10	200	1000	871.5 ± 141.6	802.4 ± 189.5	749.2 ± 171.0	740.1 ± 200.7	768.8 ± 189.2	<b>718.8 ± 169.3</b>	<b>675.2 ± 152.7</b>	765.2 ± 186.3	749.1 ± 154.2	760.8 ± 195.9	753.8 ± 178.2	747.8 ± 172.4	752.1 ± 178.2
Syn-10	200	2000	912.8 ± 174.3	930.2 ± 262.4	887.2 ± 246.1	830.6 ± 218.4	814.4 ± 200.9	<u>722.4 ± 175.8</u>	<b>707.2 ± 180.1</b>	808.1 ± 218.3	823.9 ± 243.2	863.1 ± 222.9	858.5 ± 239.2	790.2 ± 220.9	832.2 ± 231.4
<b>UK Electricity Forecasting (RMSE, kWh)</b>															
UK	100	500	2729.4 ± 8.8	2247.9 ± 51.8	2395.9 ± 59.4	2420.9 ± 45.3	2156.2 ± 51.8	<b>1990.6 ± 41.8</b>	2089.7 ± 42.2	2042.8 ± 52.1	<b>2032.3 ± 48.6</b>	2359.4 ± 54.3	2425.8 ± 55.9	2113.3 ± 43.1	2356.5 ± 46.9
UK	100	1000	2694.4 ± 9.5	1820.7 ± 38.7	2042.4 ± 36.4	1973.4 ± 39.3	1749.8 ± 29.8	<u>1634.5 ± 22.9</u>	1715.5 ± 29.4	1709.2 ± 27.4	<b>1623.2 ± 21.8</b>	2032.1 ± 37.8	1976.7 ± 34.4	1698.6 ± 26.6	1989.9 ± 39.3
UK	100	2000	2612.8 ± 7.3	1504.1 ± 19.0	1612.4 ± 16.3	1604.8 ± 16.9	1412.9 ± 10.0	<u>1354.9 ± 6.8</u>	1395.9 ± 7.9	1377.1 ± 6.9	<u>1363.1 ± 6.1</u>	1583.4 ± 18.9	1620.7 ± 19.5	1398.2 ± 10.7	1637.9 ± 19.4
<b>Image Classification (Accuracy, %)</b>															
FMNIST	200	500	90.6 ± 5.2	90.5 ± 5.2	91.5 ± 4.1	93.5 ± 3.6	91.8 ± 4.0	<b>94.5 ± 3.4</b>	91.1 ± 5.0	89.9 ± 5.2	89.9 ± 5.2	94.1 ± 3.6	<b>94.7 ± 3.2</b>	89.9 ± 5.2	92.0 ± 4.0
FMNIST	200	2000	94.5 ± 3.0	94.4 ± 3.3	94.1 ± 3.2	95.1 ± 3.0	94.2 ± 3.0	<u>95.7 ± 2.8</u>	94.4 ± 3.3	94.0 ± 3.9	95.2 ± 3.5	<b>95.8 ± 2.8</b>	94.0 ± 3.9	94.2 ± 3.2	
CIFAR	200	500	60.9 ± 4.7	64.0 ± 4.3	65.2 ± 3.9	65.5 ± 3.8	65.4 ± 3.8	<b>67.7 ± 3.7</b>	65.4 ± 3.6	64.6 ± 3.8	64.6 ± 3.8	65.5 ± 3.4	<b>67.5 ± 4.0</b>	64.6 ± 3.8	65.4 ± 3.9
CIFAR	200	2000	67.7 ± 3.9	66.7 ± 3.8	67.3 ± 3.5	67.4 ± 3.7	67.3 ± 3.4	<b>69.7 ± 3.7</b>	67.5 ± 3.6	67.0 ± 4.1	67.7 ± 3.7	<u>69.1 ± 3.6</u>	67.0 ± 4.1	67.5 ± 3.5	

- **Feature-based distances.** These distances compare task feature representations  $\mathbf{X}_v$ . We consider: (i) a normalized Euclidean distance on flattened feature matrices when shapes are consistent, or a mean-variance embedding when shapes differ (*feature distance*); (ii) a kernel-based Maximum Mean Discrepancy (MMD) with RBF kernel and median bandwidth heuristic; (iii) a Gaussian approximation using differences in feature means and covariances; and (iv) a representation-level distance based on linear Centered Kernel Alignment (CKA).
- **Target-based distances.** These distances compare task outputs  $\mathbf{y}_v$  directly. We use Euclidean distances between flattened targets (*target distance*), the symmetric KL divergence computed from smoothed histograms, the Jensen-Shannon distance, and the 1-Wasserstein (Earth Mover's) distance between empirical target distributions.
- **Optimization-geometry distances.** These distances are derived from optimization or model parameters. We consider a gradient-at-initialization proxy  $g_v = \mathbf{X}_v^\top \mathbf{y}_v$  (*gradient distance*) and a distance based on ridge-regression solutions (*model distance*), the latter serving more as a diagnostic rather than as an off-the-shelf method.

## D. Additional Experimental Results

This section provides complementary visualizations and detailed numerical results that support the main experimental findings. We report extended results for the synthetic regression benchmark and the UK electricity forecasting task experiments.

### D.1. Synthetic Regression

Figure 3 illustrates the geometry of the synthetic tasks. Increasing the within-cluster variance  $\tau_{\text{within}}$  progressively disperses task parameters while preserving the same between-cluster separation, resulting in a controlled degradation of task similarity.

Figure 4 provides a visual summary of the same sweep, highlighting the relative performance of cascaded transfer strategies across regimes.

### D.2. UK Electricity Forecasting

Figure 5 provides additional context on the UK electricity benchmark. The spatial distribution of the 100 sampled smart meters shows broad geographic coverage, while the consumption histogram highlights a strongly skewed demand distribution with a heavy right tail.

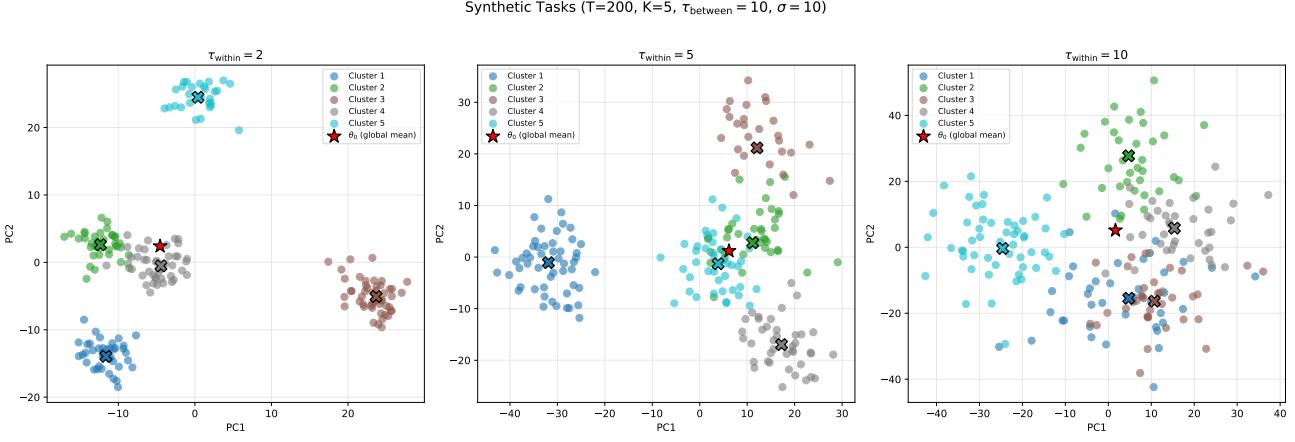


Figure 3. Two-dimensional PCA projection of task parameters for the synthetic data with increasing within-cluster variance  $\tau_{\text{within}}$ .

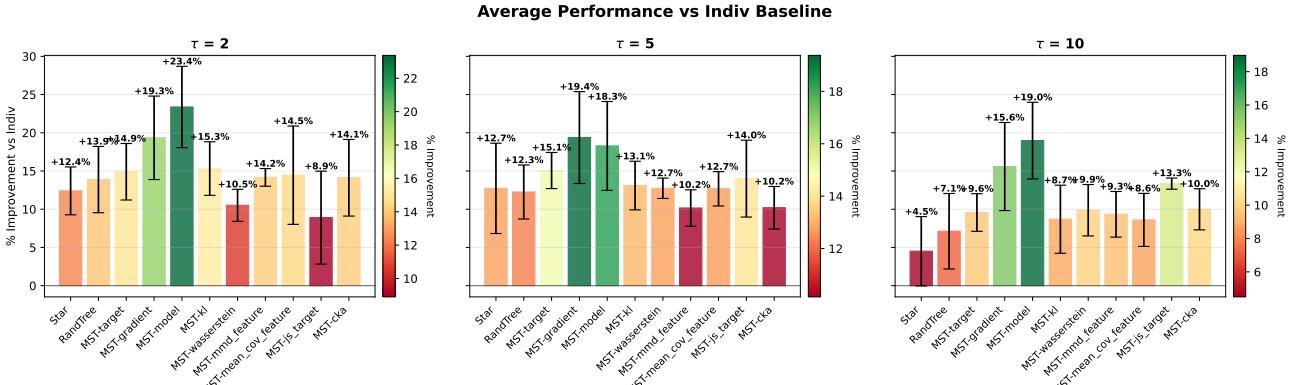
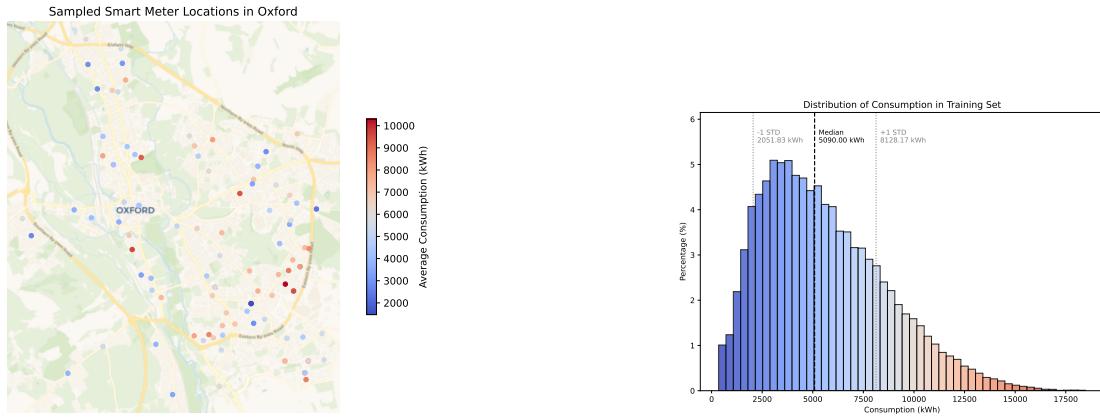


Figure 4. Test RMSE for synthetic tasks across within-cluster variability  $\tau$ , number of tasks  $T$ , and global budgets  $B$ .



(a) Spatial distribution of the sampled meters.

(b) Distribution of yearly electricity consumption.

Figure 5. Overview of spatial positions and consumption distribution of the sampled smart meters in Oxford's urban area.

## E. Budget Allocation Ablation Study

This section investigates the impact of the *budget allocation strategy* in CTL. While all main experiments use a uniform allocation of the global budget across tasks, the theoretical analysis in Section 3 suggests that allocating more refinement

effort to certain tasks (e.g., deeper nodes in the cascade or tasks connected by longer edges) may further improve performance. We empirically evaluate this question in a controlled synthetic setting.

**CTL configuration.** We evaluate CTL with minimum spanning tree construction using the gradient-based distance. The seed task is selected using the same procedure as in the main experiments, and the cascade is oriented away from the seed. Models are linear regressors trained by gradient descent. The global budget  $B$  is measured in gradient steps and is varied over  $B \in \{500, 1000, 2000\}$ . A fixed fraction (10%) of the budget is allocated to the seed task, and the remaining budget is distributed across non-root tasks according to the allocation scheme under consideration. Each task is refined exactly once.

**Budget allocation schemes.** Let  $\text{depth}(v)$  denote the depth of node  $v$  in the cascade tree, and let  $d(u, v)$  denote the edge length between a task and its parent. We compare the following allocation strategies:

- **Uniform:** all non-root tasks receive equal budget.
- **Depth-increasing:**  $b_v \propto (\text{depth}(v) + 1)^\alpha$ , allocating more budget to deeper tasks.
- **Depth-decreasing:**  $b_v \propto (\text{depth}(v) + 1)^{-\alpha}$ , favoring tasks closer to the root.
- **Edge-length-based:**  $b_v \propto (d(\text{pa}(v), v) + \varepsilon)^\beta$ , allocating more budget to tasks whose transfer edge is longer.

We consider  $\alpha \in \{1, 2\}$  and  $\beta \in \{1, 2\}$ . In all cases, real-valued allocations are converted to integer budgets using a largest-remainder scheme, ensuring that  $\sum_v b_v = B$  exactly.

**Evaluation protocol.** For each combination of  $(\tau_{\text{within}}, B)$  and allocation scheme, we repeat the experiment over 50 independent random draws of task parameters and data. Performance is measured by the RMSE, averaged across tasks. We report mean and standard deviation over repetitions. Identical trees, seeds, learning rates, and optimization procedures are used across all allocation schemes to ensure a fair comparison.

**Results.** Figure 6 reports the effect of different budget allocation strategies for CTL-MST-gradient across increasing task heterogeneity ( $\tau_{\text{within}}$ ) and total budgets  $B \in \{500, 1000, 2000\}$ . Across all regimes, uniform allocation provides a strong and reliable baseline and is never severely suboptimal. In low and moderate heterogeneity settings, uniform allocation often lies within one standard deviation of the best-performing method, particularly at larger budgets, indicating that CTL is not overly sensitive to the precise allocation rule.

At smaller budgets ( $B = 500$ ), structure-aware allocations become more influential. In several low-heterogeneity configurations, edge-length-based allocation achieves the lowest RMSE, suggesting that allocating additional refinement steps to tasks connected by longer transfer edges helps mitigate transfer bias when optimization resources are scarce. Depth-increasing allocations occasionally improve performance at intermediate budgets by compensating for accumulated upstream error at deeper nodes, whereas depth-decreasing allocations generally underperform as task heterogeneity increases, indicating that prioritizing early tasks alone is insufficient.

In the high-heterogeneity regime ( $\tau_{\text{within}} = 10$ ), performance differences between allocation schemes narrow and variance increases. While certain structure-aware schemes still achieve the best average RMSE in individual configurations, gains over uniform allocation remain modest, suggesting that budget reallocation cannot fully compensate for weak task relatedness.

Overall, these results support uniform allocation as a robust default, while confirming that simple structure-aware schemes can provide additional gains at no extra computational cost in low-budget and low-to-moderate heterogeneity regimes, in agreement with the theoretical analysis.

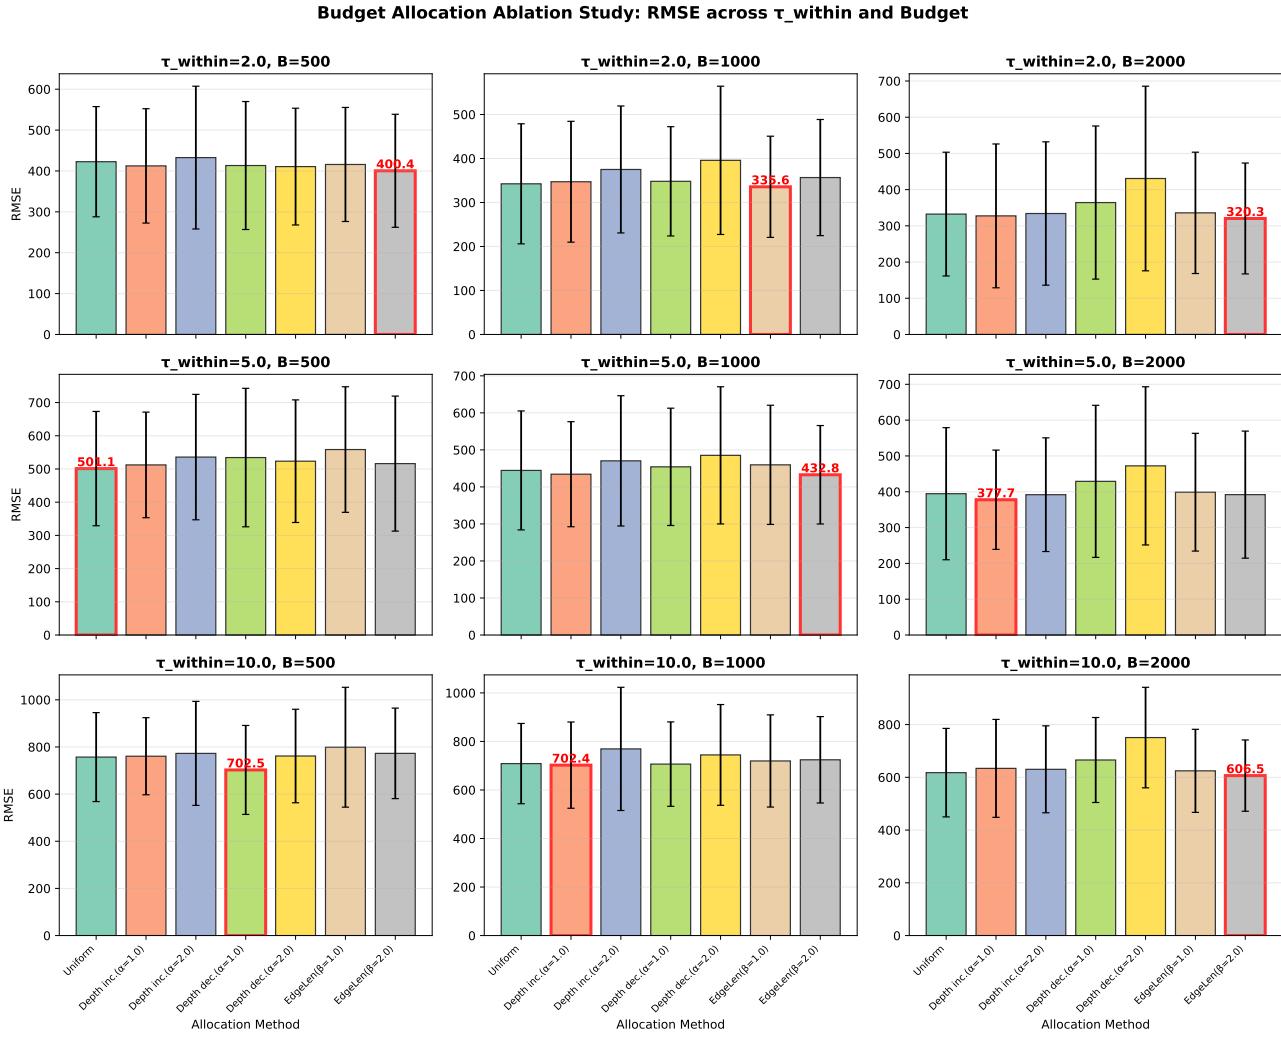


Figure 6. Budget allocation ablation for CTL-MST-gradient on clustered synthetic regression tasks. Each subplot shows mean test RMSE ( $\pm$  one standard deviation over 50 runs) for a given  $\tau_{\text{within}}$  and total budget  $B$ . Bars correspond to different budget allocation schemes, with the best-performing method highlighted.